

CIS 2107
Chapter 2 Notes
Part 2

Representing and Manipulating
Information

negative numbers?

How do we represent negative numbers?

representing negative numbers

- Sign and magnitude
- One's complement
- Two's complement
- Biased encoding


simple idea: sign and magnitude

- one bit for the sign
- everything else exactly the same
- So if this is 1_{10} :
 - 0000 0000 0000 0000 0000 0000 0000 0001₂
- -1_{10} is:
 - 1000 0000 0000 0000 0000 0000 0000 0001₂

sign and magnitude

- n bit numbers, can represent
 - -2^{n-1} to 2^{n-1}
- Some problems:
 - Two zeros:
 - $0_{10} = 0x00000000$ AND
 - $0_{10} = 0x80000000$
 - Complicated hardware

another idea: back to our odometer



odometer math

$$\begin{array}{r} 99999 \\ + 2 \\ \hline \end{array}$$

odometer math

$$\begin{array}{r} \\ \\ + \\ \hline \cancel{1} \end{array}$$

$$\begin{array}{r} \\ \\ - \\ \hline \end{array}$$

another example

$$\begin{array}{r} 9 9 9 9 9 \\ + 2 0 0 0 \\ \hline \end{array}$$

another example

$$\begin{array}{r} 9 9 9 9 \\ + 2 0 0 0 \\ \hline \cancel{1} 0 1 9 9 9 \end{array}$$

another example

$$\begin{array}{r} \\ \\ + \\ \hline \cancel{1} \\ 0 \\ 1 \\ 9 \\ 9 \\ 9 \end{array}$$

$$\begin{array}{r} \\ \\ - \\ \hline \\ 0 \\ 1 \\ 9 \\ 9 \\ 9 \end{array}$$

odometer subtraction?

- (Isn't that illegal?)
- What if we wanted to do subtraction all along?
 - subtracting 99,998 is the same as adding 2
 - subtracting 98,000 is the same as adding 2,000
- In odometer math, to subtract, what can we add?



9's complement

- decimal 9's complement of a number:
 - subtract number from all-9's (same width)
- example: 9's complement of 31,692

$$\begin{array}{r} 99,999 \\ - 31,692 \\ \hline 68,307 \end{array}$$

10's complement

- Simple. 9's complement + 1.
- 10's complement of 31,692:

$$\begin{array}{r} 99,999 \\ - 31,692 \\ \hline 68,307 \\ + 1 \\ \hline 68,308 \end{array}$$

10's complement. So What?

- In odometer math, adding 68,308 is the same as subtracting 31,692. Example:

$$\begin{array}{r} 50,000 \\ - 31,692 \\ \hline 18,308 \end{array}$$

$$\begin{array}{r} + 50,000 \\ + 68,308 \\ \hline 118,308 \end{array} \quad \rightarrow \quad \begin{array}{r} + 50,000 \\ + 68,308 \\ \hline \cancel{1}/18,308 \end{array}$$

Why this works

- losing the carry: same as subtracting 100,000

$$= 50,000 - 31,692$$

$$= 50,000 + (99,999 - 31,692 + 1) - 100,000$$

$$= 50,000 + (\cancel{99,999} - 31,692 + \cancel{1}) - \cancel{100,000}$$

one's complement

- subtract from an equal-width string of all 1's.
- example: 1's complement of 11000110_2

$$\begin{array}{r} 1111111_2 \\ - 11000110_2 \\ \hline 00111001_2 \end{array}$$

- another way of getting one's complement?

one's complement

- subtract from an equal-width string of all 1's.
- example: 1's complement of 11000110_2

$$\begin{array}{r} 1111111_2 \\ - 11000110_2 \\ \hline 00111001_2 \end{array}$$

- another way of getting one's complement?
 - bitwise NOT

one's complement arithmetic

- Addition
 - add the carry bit back in
- Subtraction
 - add the complement
- notice: like signed magnitude, two zeros

two's complement

- two's complement = one's complement + 1

two's complement example

- $6_{10} = 0110_2$
- $-6_{10}?$
- take one's complement (*i.e.*, flip bits), add 1
- $OC(0110_2) = 1001_2$
- $1001_2 + 1 = 1010_2$
- So $-6_{10} = 1010_2$

$$\text{TC}(\text{TC}(x))=x?$$

- We've said:
 - $6_{10}=0110_2$
 - $-6_{10}=1010_2$
- $\text{TC}(\text{TC}(0110_2))=0110_2?$
- $\text{TC}(1010_2)$
 - take one's complement
 - 0101_2
 - add one: 0110_2

$6_{10} = 0110_2$; $-6_{10} = 1010_2$. Check.

- So what's $6 + (-6)$?

$$\begin{array}{r} 0110_2 \\ + 1010_2 \\ \hline \cancel{1} 0000_2 \end{array}$$

TC in C

- How could we take the two's complement of a number in C?

TC in C

- How could we take the two's complement of a number in C?

```
int TC(int x) {  
    ???????  
}
```

TC in C

- How could we take the two's complement of a number in C? *Easy.*

```
int TC(int x) {  
    return -x;  
}
```

TC in C

- How could we take the two's complement of a number in C? *Easy.*

```
int TC(int x) {  
    return -x;  
}
```

- But what if we said that you couldn't use the negation operator?

TC in C without negation operator

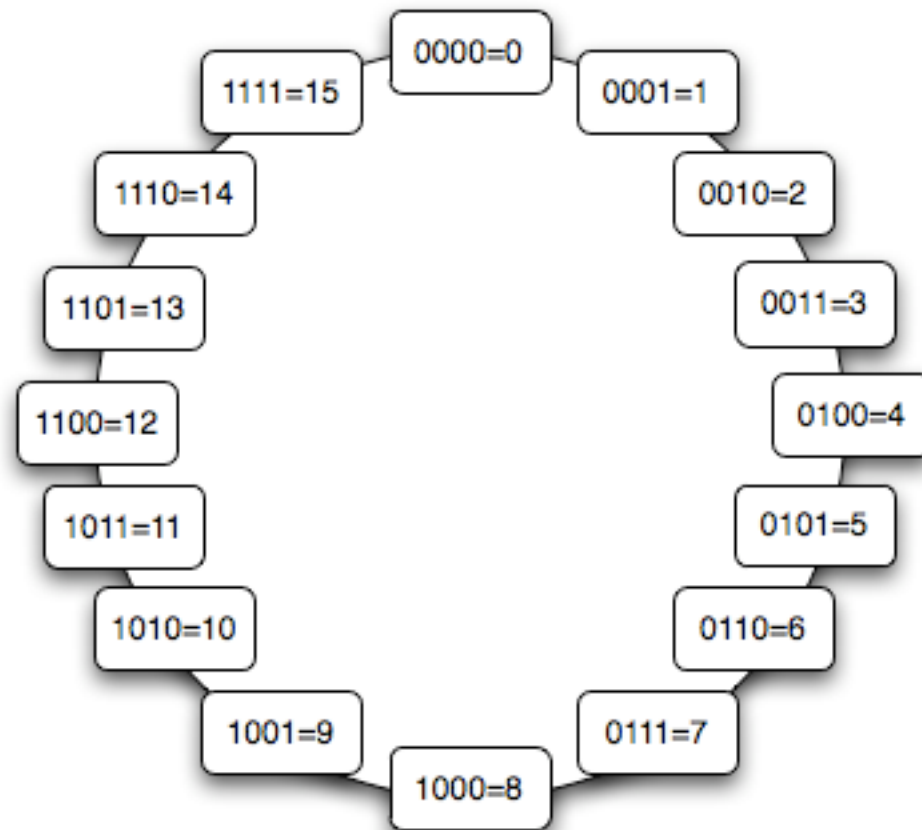
- Take the one's complement (*i.e.*, flip bits)
- Add 1
- How do we do this in C?

TC in C without negation operator

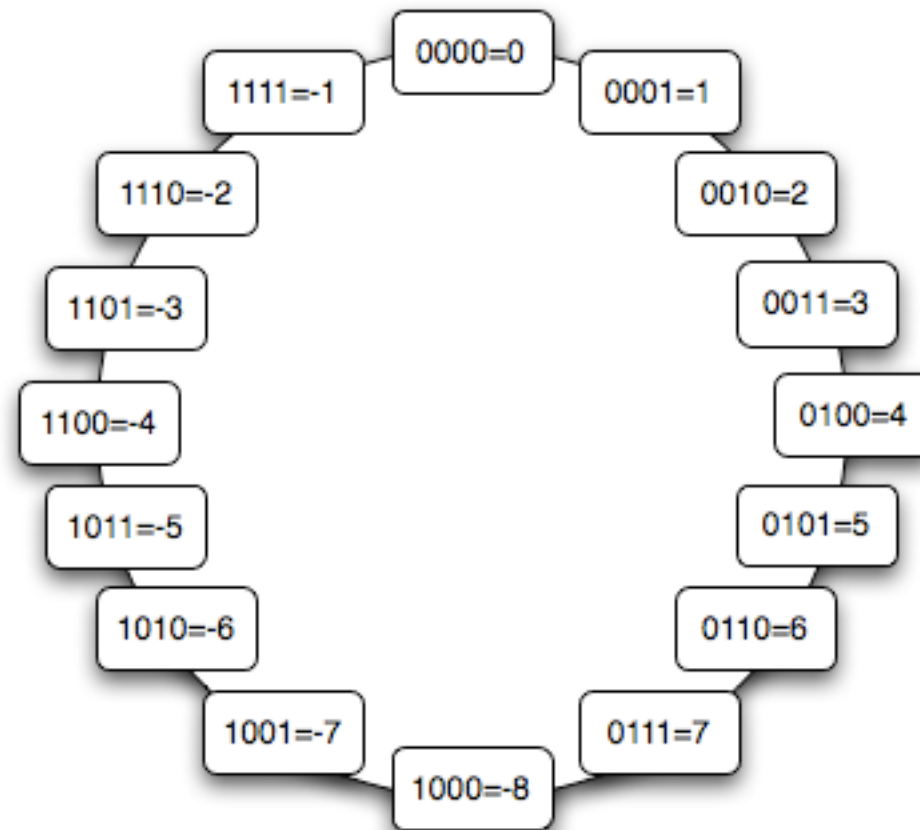
- Take the one's complement (*i.e.*, flip bits)
- Add 1
- How do we do this in C?

```
int TC(int x) {  
    return ~x+1;  
}
```

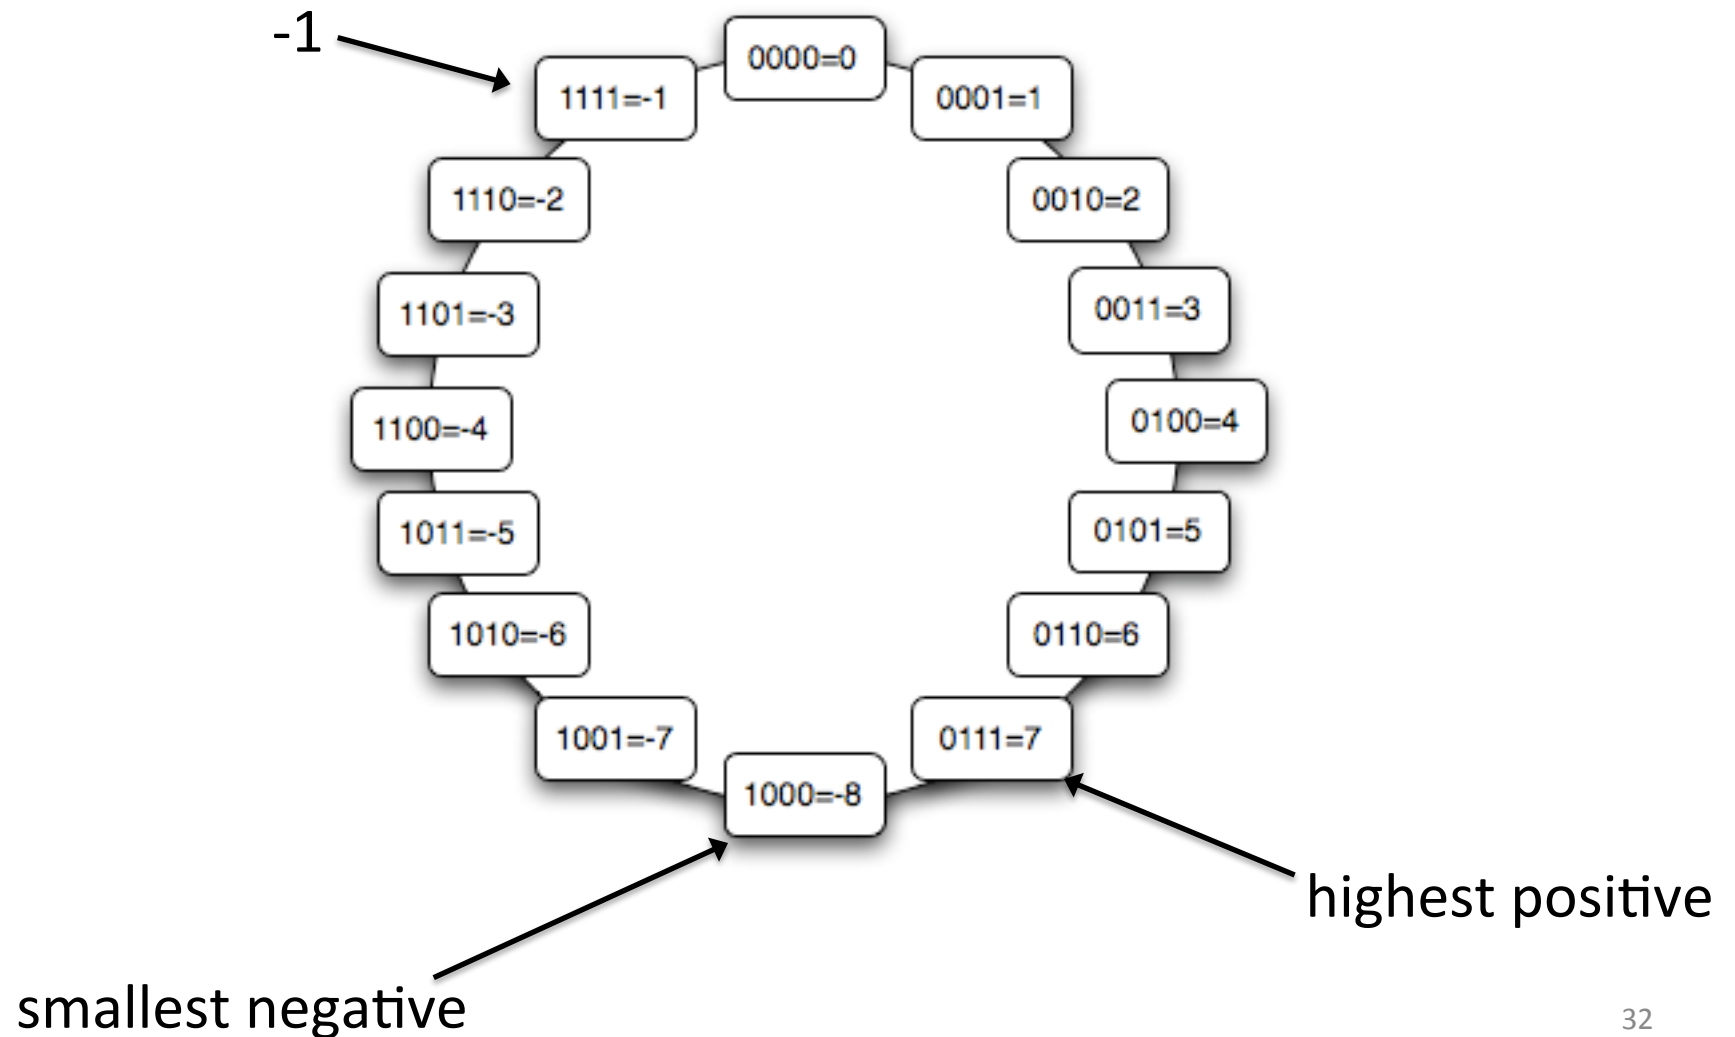
unsigned numbers



signed numbers



know the key places



key places for 16 bit words

	hex	binary
UMax	0xFFFF	0b1111111111111111
TCMax	0x7FFF	0b0111111111111111
TCMin	0x8000	0b1000000000000000
0	0x0000	0b0000000000000000
-1	0xFFFF	0b1111111111111111

key places table

word size	8	16	32	64
UMax	0xFF 255_{10}	0xFFFF $65,535_{10}$	0xFFFFFFFF $2^{32} - 1$	0xFFFFFFFFFFFFFFFF $2^{64} - 1$
TCMax	0x7F	0x7FFF	0x7FFFFFFFFF	0x7FFFFFFFFFFFFFFFFF
TCMin	0x80	0x8000	0x80000000	0x8000000000000000
0	0x00	0x0000	0x00000000	0x0000000000000000
-1	0xFF	0xFFFF	0xFFFFFFFF	0xFFFFFFFFFFFFFFFF

- know these in binary

unsigned numbers

- another way of describing unsigned nums

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

What's B2U(1010)?

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

$$\begin{aligned} B2U(1010_2) &= (1)(8) + (0)(4) + (1)(2) + (0)(1) \\ &= 10_{10} \end{aligned}$$

same for two's complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

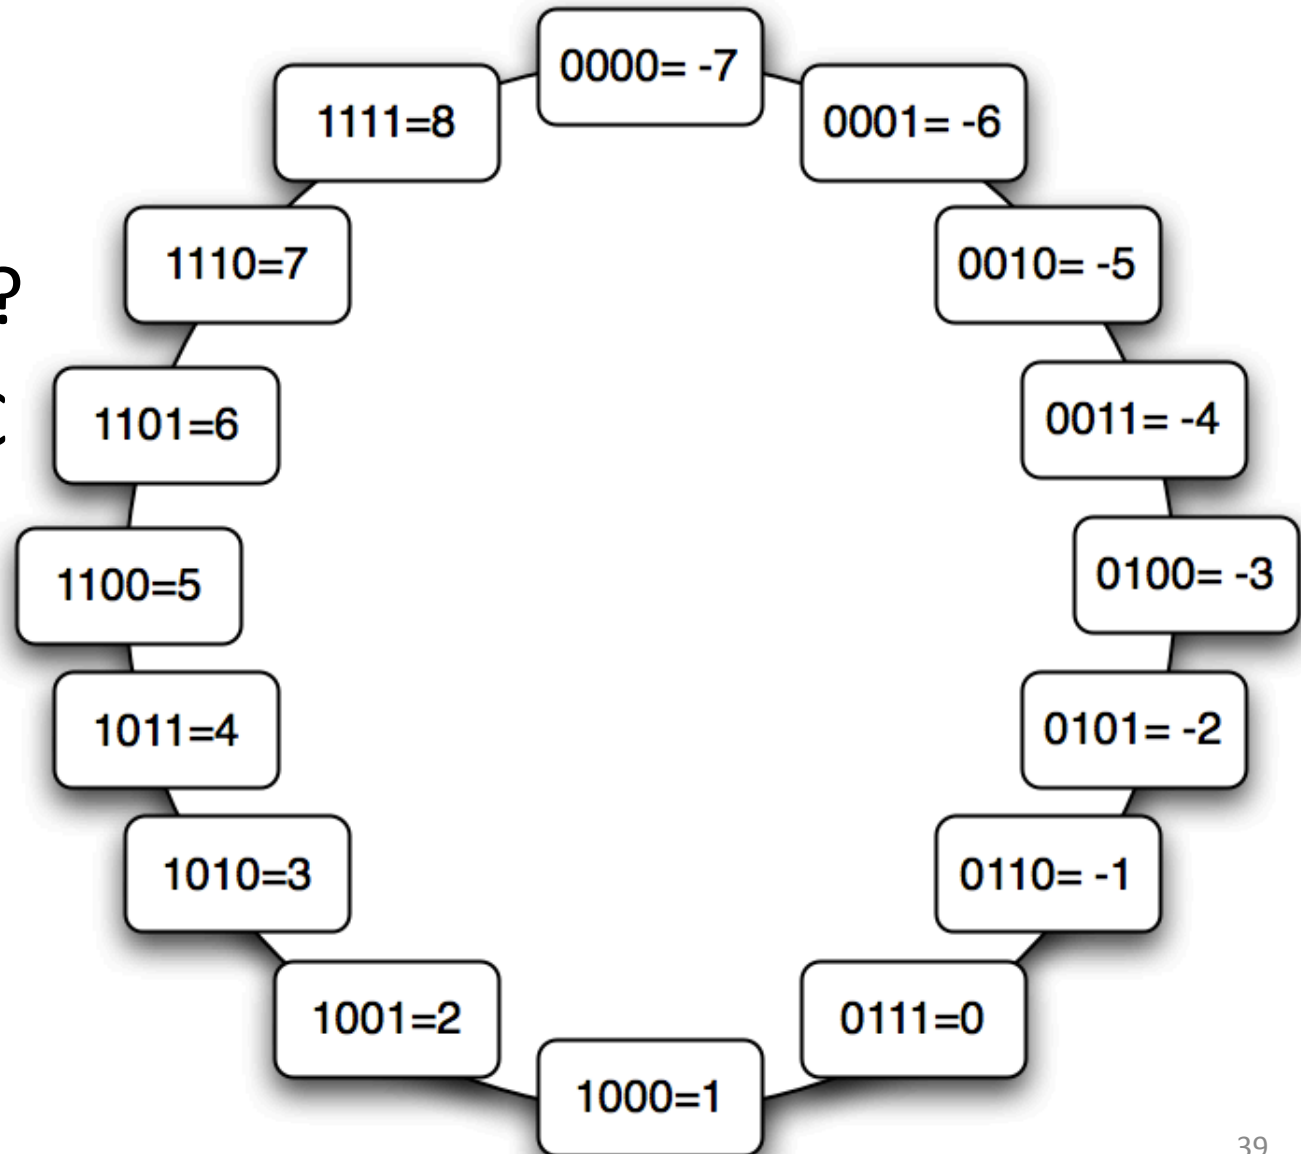
What's TC(1010)?

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

$$\begin{aligned} B2T(1010) &= (1)(\mathbf{-8}) + (0)(4) + (1)(2) + (0)(1) \\ &= -6 \end{aligned}$$

Bias Encoding. Bias= $-2^{\text{width}-1}-1$

- # zeros?
- # positives?
- order vs TC order



binary	unsigned	TC	OC	SM
0000	0	0	0	0
0001	1	1	1	1
0010	2	2	2	2
0011	3	3	3	3
0100	4	4	4	4
0101	5	5	5	5
0110	6	6	6	6
0111	7	7	7	7
1000	8	-8	-7	-0
1001	9	-7	-6	-1
1010	10	-6	-5	-2
1011	11	-5	-4	-3
1100	12	-4	-3	-4
1101	13	-3	-2	-5
1110	14	-2	-1	-6
1111	15	-1	-0	-7

C integer types

- char, short int, int, long int (C99 long long int)
- unsigned versions of each
- C spec doesn't say, but most implementations TC

How big are each?

C data type	typical 32-bit	Intel IA32	x86-64
char	1	1	1
short	2	2	2
int	4	4	4
long	4	4	8
long long	8	8	8
float	4	4	4
double	8	8	8
long double	8	10/12	10/16
pointer	4	4	8

Is this sort of thing an issue in Java?

constants

integers

- 123 (32-bit int)
- 123L (32-bit long int)
- 123u (32-bit unsigned int),
- 123uL (32-bit unsigned long)
- 123LL (64-bit signed long long)

octal and hex

- 0100 (100 octal = 64 decimal)
- 0x100 (100 hex = 256 decimal)
- 0xful (15 unsigned long).

floating point

- 12.3 (32-bit float),
- 123e-1 (32-bit float)
- 12.3f (32-bit float),
- 12.3L (64-bit long double)

'x' = character constant (0 to 127).

- In ASCII ' ' = '\040' = 'x20' = 32
- In ASCII '0' = '\060' = 'x30' = 48
- In ASCII 'A' = '\101' = 'x41' = 65
- In ASCII 'a' = '\141' = 'x61' = 97

What range can be stored?

width	signed?	range
8 bits	unsigned	$0 \text{ to } 2^8 - 1$
	signed	$-2^7 \text{ to } 2^7 - 1$
16 bits	unsigned	$0 \text{ to } 2^{16} - 1$
	signed	$-2^{15} \text{ to } 2^{15} - 1$
32 bits	unsigned	$0 \text{ to } 2^{32} - 1$
	signed	$-2^{31} \text{ to } 2^{31} - 1$
64 bits	unsigned	$0 \text{ to } 2^{64} - 1$
	signed	$-2^{63} \text{ to } 2^{63} - 1$

limits.h

```
/* Number of bits in a 'char'. */
# define CHAR_BIT      8

/* Minimum and maximum values a 'signed char' can hold. */
# define SCHAR_MIN      (-128)
# define SCHAR_MAX      127

/* Maximum value an 'unsigned char' can hold. (Minimum is 0.) */
# define UCHAR_MAX      255

/* Minimum and maximum values a 'signed short int' can hold. */
# define SHRT_MIN        (-32768)
# define SHRT_MAX        32767

/* Maximum value an 'unsigned short int' can hold. (Minimum is 0.) */
# define USHRT_MAX      65535

/* Minimum and maximum values a 'signed int' can hold. */
# define INT_MIN          (-INT_MAX - 1)
# define INT_MAX          2147483647

/* Maximum value an 'unsigned int' can hold. (Minimum is 0.) */
# define UINT_MAX        4294967295U

/* Minimum and maximum values a 'signed long int' can hold. */
# define LONG_MAX        2147483647L
# define LONG_MIN        (-LONG_MAX - 1L)
```

C99 <stdint.h>

- for all widths W that the machine supports
 - exact width types: `intW_t`, `uintW_t`
 - *e.g.*, `int8_t`, `uint32_t`.
 - minimum width types: `int_leastW_t`, `uint_leastW_t`
 - <limits.h>-style macros for these types:
 - *e.g.* `INTW_MIN`, `UINTW_MAX`, ...

take a look at some of these with gdb

```
int p1 = 37;  
unsigned int p2 = 37;  
int n1 = -37;
```

- remember to compile:
 - with the `-g` switch
 - don't turn on optimization (*no* `-O`)
- in gdb, to print binary, use `p/t`
- confirm you get the expected when you try:
 - `p/d (~n1+1)`

What happens?

```
int x=-1;  
unsigned int ux = (unsigned int) x;
```


What happens?

```
int x=-1;
unsigned int ux = (unsigned int) x;
```


```
(gdb) p/t x  
$1 = 1111111111111111111111111111111111111111111111111
```

```
(gdb) p/d x
$2 = -1
```

```
(gdb) p/t ux  
$3 = 11111111111111111111111111111111111111111111111
```

```
(gdb) p/u ux
$4 = 4294967295
```

casting from signed to unsigned



what happens if we cast a negative to an equal width unsigned?

bits
0000 ₂
0001 ₂
0010 ₂
0011 ₂
0100 ₂
0101 ₂
0110 ₂
0111 ₂
1000 ₂
1001 ₂
1010 ₂
1011 ₂
1100 ₂
1101 ₂
1110 ₂
1111 ₂


signed
00 ₁₀
01 ₁₀
02 ₁₀
03 ₁₀
04 ₁₀
05 ₁₀
06 ₁₀
07 ₁₀
−08 ₁₀
−07 ₁₀
−06 ₁₀
−05 ₁₀
−04 ₁₀
−03 ₁₀
−02 ₁₀
−01 ₁₀

unsigned
00 ₁₀
01 ₁₀
02 ₁₀
03 ₁₀
04 ₁₀
05 ₁₀
06 ₁₀
07 ₁₀
08 ₁₀
09 ₁₀
10 ₁₀
11 ₁₀
12 ₁₀
13 ₁₀
14 ₁₀
15 ₁₀

bits
0000 ₂
0001 ₂
0010 ₂
0011 ₂
0100 ₂
0101 ₂
0110 ₂
0111 ₂
1000 ₂
1001 ₂
1010 ₂
1011 ₂
1100 ₂
1101 ₂
1110 ₂
1111 ₂

signed		unsigned
00 ₁₀	same	00 ₁₀
01 ₁₀		01 ₁₀
02 ₁₀		02 ₁₀
03 ₁₀		03 ₁₀
04 ₁₀		04 ₁₀
05 ₁₀		05 ₁₀
06 ₁₀		06 ₁₀
07 ₁₀		07 ₁₀
−08 ₁₀	+16	08 ₁₀
−07 ₁₀		09 ₁₀
−06 ₁₀		10 ₁₀
−05 ₁₀		11 ₁₀
−04 ₁₀		12 ₁₀
−03 ₁₀		13 ₁₀
−02 ₁₀		14 ₁₀
−01 ₁₀		15 ₁₀

casting from unsigned to signed



what if we cast a large unsigned positive to an equal width signed?

bit shifting for unsigned numbers

- \gg , \ll , $\gg=$, $\ll=$

0	0	1	0	1	1	0	1
---	---	---	---	---	---	---	---

- numbers fall off the end

shift left 1

0	1	0	1	1	0	1	0
---	---	---	---	---	---	---	---

- fill in with 0s

- math equivalent (when 1s haven't fallen off)?

shift right 2

0	0	0	1	0	1	1	0
---	---	---	---	---	---	---	---

shifting left

- $x \ll j$
 - shift x to the left j bit positions
 - fill with 0s from the right
 - numbers “fall off the left end”

```
char x=11, j;
```

```
for (j=0; j<8; j++)  
    x<<=1;
```

	x_{10}	x_2
0	11	0b00001011
1	22	0b00010110
2	44	0b00101100
3	88	0b01011000
4	176	0b10110000
5	96	0b01100000
6	192	0b11000000
7	128	0b10000000

shifting to the right

- unsigned:
 - same as left shift
- what about signed?
 - implementation dependent
 - some fill from LHS
 - with 0
 - with sign bit
 - (why do this with signed numbers anyway?)

```
int main(void)
{
    int i;
    char c1 = 64, c2 = -64;
    for (i=8; i>0; i--) {
        c1>>=1;
        c2>>=1;
    }
    return 0;
}
```


(aside) java right shift

- Java defines two:
 - `>>` fills from the left with the sign bit
 - `>>>` fills from the left with 0s

some example code

```
1 public class JavaShift {
2     public static void main(String args[]) {
3         int x = Integer.MIN_VALUE; // i.e. -2**(31)
4         for (int i=0; i<32; i++) {
5             int ds = x>>i;
6             int ts = x>>>i;
7             System.out.println(x+" >> "+ i + " = " +
8                                 Integer.toBinaryString(ds));
9             System.out.println(x+" >>> "+ i + " = " +
10                                Integer.toBinaryString(ts));
11             System.out.println();
12         }
13     }
14 }
```

java right shifts output

```
-2147483648 >> 0 = 10000000000000000000000000000000  
-2147483648 >>> 0 = 10000000000000000000000000000000  
  
-2147483648 >> 1 = 11000000000000000000000000000000  
-2147483648 >>> 1 = 10000000000000000000000000000000  
  
-2147483648 >> 2 = 11100000000000000000000000000000  
-2147483648 >>> 2 = 10000000000000000000000000000000  
  
-2147483648 >> 3 = 11110000000000000000000000000000  
-2147483648 >>> 3 = 10000000000000000000000000000000  
  
-2147483648 >> 4 = 11111000000000000000000000000000  
-2147483648 >>> 4 = 10000000000000000000000000000000  
  
-2147483648 >> 5 = 11111100000000000000000000000000  
-2147483648 >>> 5 = 10000000000000000000000000000000
```

casting to different widths

- smaller to larger
 - no problem
- larger to smaller
 - information loss?

small to large

- What happens?
 - long x = char c

unsigned: zero extension

- unsigned numbers
 - going from small to large → zero extension
- Example: unsigned long x = unsigned char c
- What happens?

signed: sign extension

- signed numbers
 - going from small to large \rightarrow sign extension
- Example: `long x = char c`
- What happens?
 - if `c > 0`?
 - if `c < 0`?

C conversion rules

- Details in K&R Appendix A
- C rules. Three types:
 - Integer promotion
 - Integer conversion rank
 - Usual arithmetic conversions

C Conversion Rules: Integer Promotion

- Integer types smaller than int promoted to int.
- example:
 - char result, a=100, b=10, c=20

C Conversion Rules: Integer Conversion Rank

C Conversion Rules: “Usual” Arithmetic Conversions

casting large to small. truncation

```
int x1 = 0x00001234;  
int x2 = 0x12345678;  
short s1 = (short)x1;  
short s2 = (short)x2;  
printf("x1=0x%08x, ", x1);  
printf("x2=0x%08x, ", x2);  
printf("s1=0x%08x, ", s1);  
printf("s2=0x%08x\n", s2);
```

we get:

x1=0x00001234, x2=0x12345678, s1=0x00001234, s2=0x00005678

truncation of unsigned numbers

- unsigned int x
- truncating to k bits equivalent to $x \bmod 2^k$

integer arithmetic

- “odometer” effect
- modular arithmetic
- unsigned addition: addition modulo width
- for others, see details in BO

unsigned addition

- addition modulo the width

@ @ @ TAKE A LOOK AT BO SLIDES 28

--@ @ @

@ @ @ ADD CMU SL 31 MATERIAL
@ @ @

Floating Point

What happens here?

```
1  #include <stdio.h>
2
3  int main(int argc, char **argv)
4  {
5      float f=0.1;
6
7      if (f==0.1)
8          printf("It's 0.1\n");
9      else
10         printf("It's not 0.1\n");
11
12     return 0;
13 }
```

What happens here?

```
1  #include <stdio.h>
2
3  int main(int argc, char **argv)
4  {
5      float f=0.1;
6
7      if (f==0.1)
8          printf("It's 0.1\n");
9      else
10         printf("It's not 0.1\n");
11
12     return 0;
13 }
```

```
1  bash-3.2$ gcc -Wall -o StrangeFloat02 StrangeFloat02.c
2  bash-3.2$ ./StrangeFloat02
3  It's not 0.1
```

What about here?

```
1  #include <stdio.h>
2
3  int main(int argc, char **argv)
4  {
5      float f1=0.1,
6          f2=0.2,
7          sum;
8
9      sum=f1+f2;
10     printf("%.8f+%.8f=%.8f\n", f1, f2, sum);
11
12     return 0;
13 }
```

What about here?

```
1  #include <stdio.h>
2
3  int main(int argc, char **argv)
4  {
5      float f1=0.1,
6           f2=0.2,
7           sum;
8
9      sum=f1+f2;
10     printf("%.8f+%.8f=%.8f\n", f1, f2, sum);
11
12     return 0;
13 }
```

```
1  bash-3.2$ gcc -Wall -o StrangeFloat StrangeFloat.c
2  bash-3.2$ ./StrangeFloat
3  0.10000000+0.20000000=0.30000001
```

Strange on 32-bit machines

```
1  #include <stdio.h>
2
3  int main(void) {
4      float a = 300000000;
5      float b = 3;
6      float c;
7
8      c = a + b - a;
9      printf("%f\n", c);
10     c = a + b;
11     c = c - a;
12     printf("%f\n", c);
13
14     return 0;
15 }
```

Strange on 32-bit machines

```
1  #include <stdio.h>
2
3  int main(void) {
4      float a = 300000000;
5      float b = 3;
6      float c;
7
8      c = a + b - a;
9      printf("%f\n", c);
10     c = a + b;
11     c = c - a;
12     printf("%f\n", c);
13
14     return 0;
15 }
```

Output:

```
3.000000
4.000000
```


Recall Binary Representation of ints

$$11011001_2 =$$

2^7	2^6	2^5	2^4	2^3	2^2	2^1	2^0
1	1	0	1	1	0	0	1

$$= (1)(128) + (1)(64) + (0)(32) + (1)(16) + (1)(8) + (0)(4) + (0)(2) + (1)(1)$$

$$= 128 + 64 + 16 + 8 + 1$$

$$= 217$$

What about fractions?

$$.11011001_2 =$$

.	2^{-1}	2^{-2}	2^{-3}	2^{-4}	2^{-5}	2^{-6}	2^{-7}	2^{-8}
.	1	1	0	1	1	0	0	1

$$= 2^{-1} + 2^{-2} + 2^{-4} + 2^{-5} + 2^{-8}$$

$$= 0.5 + 0.25 + 0.0625 + 0.03125 + 0.00390625$$

$$= 0.84765625$$

some powers of 2 ≤ 1

2^0	1.000000
2^{-1}	0.500000
2^{-2}	0.250000
2^{-3}	0.125000
2^{-4}	0.062500
2^{-5}	0.031250
2^{-6}	0.015625
2^{-7}	0.0078125
2^{-8}	0.00390625

Decimal to binary?

- How do we convert the representation of:
 - a fraction in decimal
 - to a fraction in binary

recall how we did numbers > 1

- Convert 119_{10} to binary
 - solution: $111\ 0111_2$
- $$\begin{aligned}119 &= 59 * 2 + 1 \\59 &= 29 * 2 + 1 \\29 &= 14 * 2 + 1 \\14 &= 7 * 2 + 0 \\7 &= 3 * 2 + 1 \\3 &= 1 * 2 + 1 \\1 &= 0 * 2 + 1\end{aligned}$$

converting fractions

```
1  while (fraction part != 0) {  
2      multiply number by 2  
3      record the integer part for later  
4      subtract the integer part  
5  }  
6  
7  recorded integer parts are the binary rep
```

example. 0.6953125_{10} ?

$$0.6953125_{10} * 2 = 1.390625_{10}$$

$$0.390625_{10} * 2 = \mathbf{0.78125_{10}}$$

$$0.78125_{10} * 2 = 1.5625_{10}$$

$$0.5625_{10} * 2 = 1.125_{10}$$

$$0.125_{10} * 2 = \mathbf{0.25_{10}}$$

$$0.25_{10} * 2 = \mathbf{0.5_{10}}$$

$$0.5_{10} * 2 = 1.0_{10}$$

solution: $0.6953125_{10} = 0.1011001_2$

double check result

- $0.1011001_2 = 0.6953125_{10}????$

$$0.1011001_2 = 2^{-1} + 2^{-3} + 2^{-4} + 2^{-7}$$

$$= 0.5_{10} + 0.125_{10} + 0.0625_{10} + 0.0078125_{10}$$

$$\begin{array}{r} 0.5_{10} \\ 0.125_{10} \\ 0.0625_{10} \\ + \quad 0.0078125_{10} \\ \hline 0.6953125_{10} \end{array}$$

Does it always work?

- Remember the pigeonhole?
 - infinite number of floating-point numbers
 - storing in register of finite size

Does it always work? Repeating

- We have repeating decimal fractions
- We also have repeating binary fractions.
 - $0.1_{10} = 0.0001100110011_2 \dots$
 - $0.2_{10} = 0.001100110011_2 \dots$
- The more places we have, the closer we are to the value we're trying to represent

representing 0.2_{10}

base 2 float	base 10 frac	base 10 float
0.0_2	0_{10}	0_{10}
0.00_2	0_{10}	0_{10}
0.001_2	$1/8_{10}$	0.125_{10}
0.0011_2	$3/16_{10}$	0.1875_{10}
0.00110_2	$3/16_{10}$	0.1875_{10}
0.001100_2	$3/16_{10}$	0.1875_{10}
0.0011001_2	$25/128_{10}$	0.1953125_{10}
0.00110011_2	$51/256_{10}$	0.19921875_{10}
0.001100110_2	$51/256_{10}$	0.19921875_{10}
0.0011001100_2	$51/256_{10}$	0.19921875_{10}
0.00110011001_2	$409/2048_{10}$	0.19970703125_{10}
0.001100110011_2	$819/4096_{10}$	0.199951171875_{10}
0.0011001100110_2	$819/4096_{10}$	0.199951171875_{10}
0.00110011001100_2	$819/4096_{10}$	0.199951171875_{10}
0.001100110011001_2	$6553/32768_{10}$	0.199981689453125_{10}
0.0011001100110011_2	$13107/65536_{10}$	0.1999969482421875_{10}
0.00110011001100110_2	$13107/65536_{10}$	0.1999969482421875_{10}
0.001100110011001100_2	$13107/65536_{10}$	0.1999969482421875_{10}
0.0011001100110011001_2	$104857/524288_{10}$	$0.1999988555908203125_{10}$
0.00110011001100110011_2	$209715/1048576_{10}$	$0.19999980926513671875_{10}^{91}$

representing 0.1_{10}

base 2 float	base 10 frac	base 10 float
0.00_2	0_{10}	0_{10}
0.000_2	0_{10}	0_{10}
0.0001_2	$1/16_{10}$	0.0625_{10}
0.00011_2	$3/32_{10}$	0.09375_{10}
0.000110_2	$3/32_{10}$	0.09375_{10}
0.0001100_2	$3/32_{10}$	0.09375_{10}
0.00011001_2	$25/256_{10}$	0.09765625_{10}
0.000110011_2	$51/512_{10}$	0.099609375_{10}
0.0001100110_2	$51/512_{10}$	0.099609375_{10}
0.00011001100_2	$51/512_{10}$	0.099609375_{10}
0.000110011001_2	$409/4096_{10}$	0.099853515625_{10}
0.0001100110011_2	$819/8192_{10}$	0.0999755859375_{10}
0.00011001100110_2	$819/8192_{10}$	0.0999755859375_{10}
0.000110011001100_2	$819/8192_{10}$	0.0999755859375_{10}
0.0001100110011001_2	$6553/65536_{10}$	0.0999908447265625_{10}
0.00011001100110011_2	$13107/131072_{10}$	0.09999847412109375_{10}
0.000110011001100110_2	$13107/131072_{10}$	0.09999847412109375_{10}
0.0001100110011001100_2	$13107/131072_{10}$	0.09999847412109375_{10}
0.00011001100110011001_2	$104857/1048576_{10}$	$0.09999942779541015625_{10}$
$0.000110011001100110011_2$	$209715/2097152_{10}$	$0.099999904632568359375_{10}$

can't represent everything

- with finite length binary strings, can only approximate numbers that can't be written as:

$$(x)(2^y)$$

- (remember that y can be positive or negative)

Other “interesting” numbers

base 2 float	base 10 frac	base 10 float
0.1_2	$1/2_{10}$	0.5_{10}
0.11_2	$3/4_{10}$	0.75_{10}
0.111_2	$7/8_{10}$	0.875_{10}
0.1111_2	$15/16_{10}$	0.9375_{10}
0.11111_2	$31/32_{10}$	0.96875_{10}
0.111111_2	$63/64_{10}$	0.984375_{10}
0.1111111_2	$127/128_{10}$	0.9921875_{10}
0.11111111_2	$255/256_{10}$	0.99609375_{10}
0.111111111_2	$511/512_{10}$	0.998046875_{10}
0.1111111111_2	$1023/1024_{10}$	0.9990234375_{10}

machine representation of floats

- represent as:

$$\pm m \times b^e$$

m mantissa

b base

e exponent

- but, base is always 2

machine representation of floats

- Recall: 1234567.0_{10} can be written:
 - $123456.7 * 10$
 - $12345.67 * 10^2$
 - $1234.567 * 10^3$
 - $123.4567 * 10^4$
 - $12.34567 * 10^5$
 - $1.234567 * 10^6$

machine representation of floats

- 110010.0010_2 can be written:
 - $110010.0010_2 * 2^0$
 - $11001.00010_2 * 2^1$
 - $1100.100010_2 * 2^2$
 - $110.0100010_2 * 2^3$
 - $11.00100010_2 * 2^4$
 - $1.100100010_2 * 2^5$
- Shift until radix after first 1 called *normalizing*

machine representation of floats



floating point representation



C type	exponent	mantissa
float	8 bits	23 bits
double	11 bits	52 bits

floating point representation

sign	exponent	mantissa
------	----------	----------

C type	exponent	mantissa
float	8 bits	23 bits
double	11 bits	52 bits

On Intel – “extended precision” (depending on compiler)

C type	exponent	mantissa
long double	15 bits	64 bits

How?

- How are floating-point numbers represented?
- There are three cases.

--- let's do the most common case first

How? The common case.

- Set the sign bit. 0 for positive, 1 for negative
- Write num in fixed-pt binary
- Normalize (radix pt is just to the right of the first 1)
- m is the values to the right of radix point
- calculate e : $\text{exponent} + \text{bias}$
 - $\text{bias} = 2^{\text{exponent field width}-1} - 1$
 - for floats, bias is $2^{8-1} - 1 = 2^7 - 1 = 127$
 - for doubles, it's $2^{11-1} - 1 = 2^{10} - 1 = 1023$
 - Can represent exponents of:
 - -126 to +127 for floats
 - -1022 to +1023 for doubles

How? Example -15.375

- Set the sign bit. 0 for positive, 1 for negative
 - sign bit 1
- Write num in fixed-pt binary:
 - $15.375 = 1111.011_2$
- Normalize (so radix point is just to the right of the first 1)
 - $1.111011_2 * 2^3$
- m is the values to the right of radix point
 - 111011
- calculate e : exponent+127
 - exponent was 3_{10} (11_2).
 - $3_{10} + 127_{10} = 130_{10}$ or $1000\ 0010_2$
- Final result:

1

1000 0010


111 0110 0000 0000 0000 0000

Another example

- -118.625

Another example

- -118.625



double check: try in gdb


- -118.625?

1	1000 0101	110 1101 0100 0000 0000 0000
---	-----------	------------------------------
- Check.
- In a program where we have:
 - `float fl`
- `(gdb) set fl=-118.625`
- `(gdb) x/t &fl`
- `0x7fff5fbff68c:`
`11000010111011010100000000000000`

How?

- How are floating-point numbers represented?
- There are three cases:
 - *Normalized* values (what we just did)
 - *De-normalized* values
 - Numbers “close” to 0.0
 - Exponent field is all zeros
 - “Special” cases
 - $\pm \infty$, NaN
 - Exponent field is all 1s.

Visualizing Floating-point Range



How?

- How are floating-point numbers represented?
- There are three cases:
 - *Normalized* values (what we just did)
 - ***De-normalized* values**
 - Numbers “close” to 0.0
 - Exponent field is all zeros
 - “Special” cases
 - $\pm \infty$, NaN
 - Exponent field is all 1s.

De-Normalized Values

- Why?
 - Used to represent values “close” to 0.0
- *Exponent field* – all 0s.
 - Fraction represented has exponent of 1-Bias
- *Mantissa field*
 - we don’t assume a leading 1
- *Sign field*
 - As usual, can be 1 or 0.
 - Means that we can also have +0.0 or -0.0

Example

- We'll use 7-bit floats, consisting of:
 - A sign bit
 - 3 bits for the exponent
 - 3 bits for the mantissa
- What is: 0 101 000, where:
 - 0 is the sign bit
 - 101 are the bits for the mantissa
 - 000 are the bits for the exponent

Example: 0 101 000

- We'll use 7-bit floats, consisting of:
 - A sign bit
 - 3 bits for the exponent
 - 3 bits for the mantissa
- What is: 0 101 000, where:
 - 0 is the sign bit
 - 101 are the bits for the mantissa
 - 000 are the bits for the exponent
- Bias is $2^{(3-1)} - 1 = 2^2 - 1 = 3$
- Exponent is $1 - \text{bias} = -2$
- Mantissa:
 - $\frac{1}{2} + 0/4 + 1/8 = 5/8$
- Final result:
 - Mantissa * 2^{exponent} =
 - $5/8 * 2^{-2} =$
 - $5/8 * 1/4 =$
 - $5/32 =$
 - 0.15625

Another de-normalized example

- 0 110 000 (110 is the mantissa)

Another de-normalized example

- 0 110 000 (110 is the mantissa)
- Bias is $2^{(3-1)}-1 = 2^2-1 = 3$
- Exponent is $1-\text{bias} = -2$
- $M = \frac{1}{2} + \frac{1}{4} + \frac{0}{8} = \frac{3}{4}$
- Final result = $M * 2^{\text{exponent}} =$
 - $\frac{3}{4} * 2^{-2} =$
 - $\frac{3}{4} * \frac{1}{4} =$
 - $\frac{3}{16} = 0.1875$

How?

- How are floating-point numbers represented?
- There are three cases:
 - *Normalized* values (what we just did)
 - *De-normalized* values
 - Numbers “close” to 0.0
 - Exponent field is all zeros
 - **“Special” cases**
 - $\pm \infty$, NaN
 - Exponent field is all 1s.

“Special” Cases

- The exponent field is all 1s.
- If the fraction field is all 0s:
 - +/- infinity
 - Can use +/- infinity when:
 - overflow has occurred
 - divide by 0.
- Otherwise:
 - NaN, *e.g.* `sqrt(-1)`

“Simple” examples

- 32-, 64-, or 80-bit widths: tough to see
- keep it simple for now: 8-bit widths.
 - 1 bit for sign
 - 4 bits for the exponent
 - 3 bits for the fraction

Bias

- With a 4-bit exponent, what will be the bias?

Bias

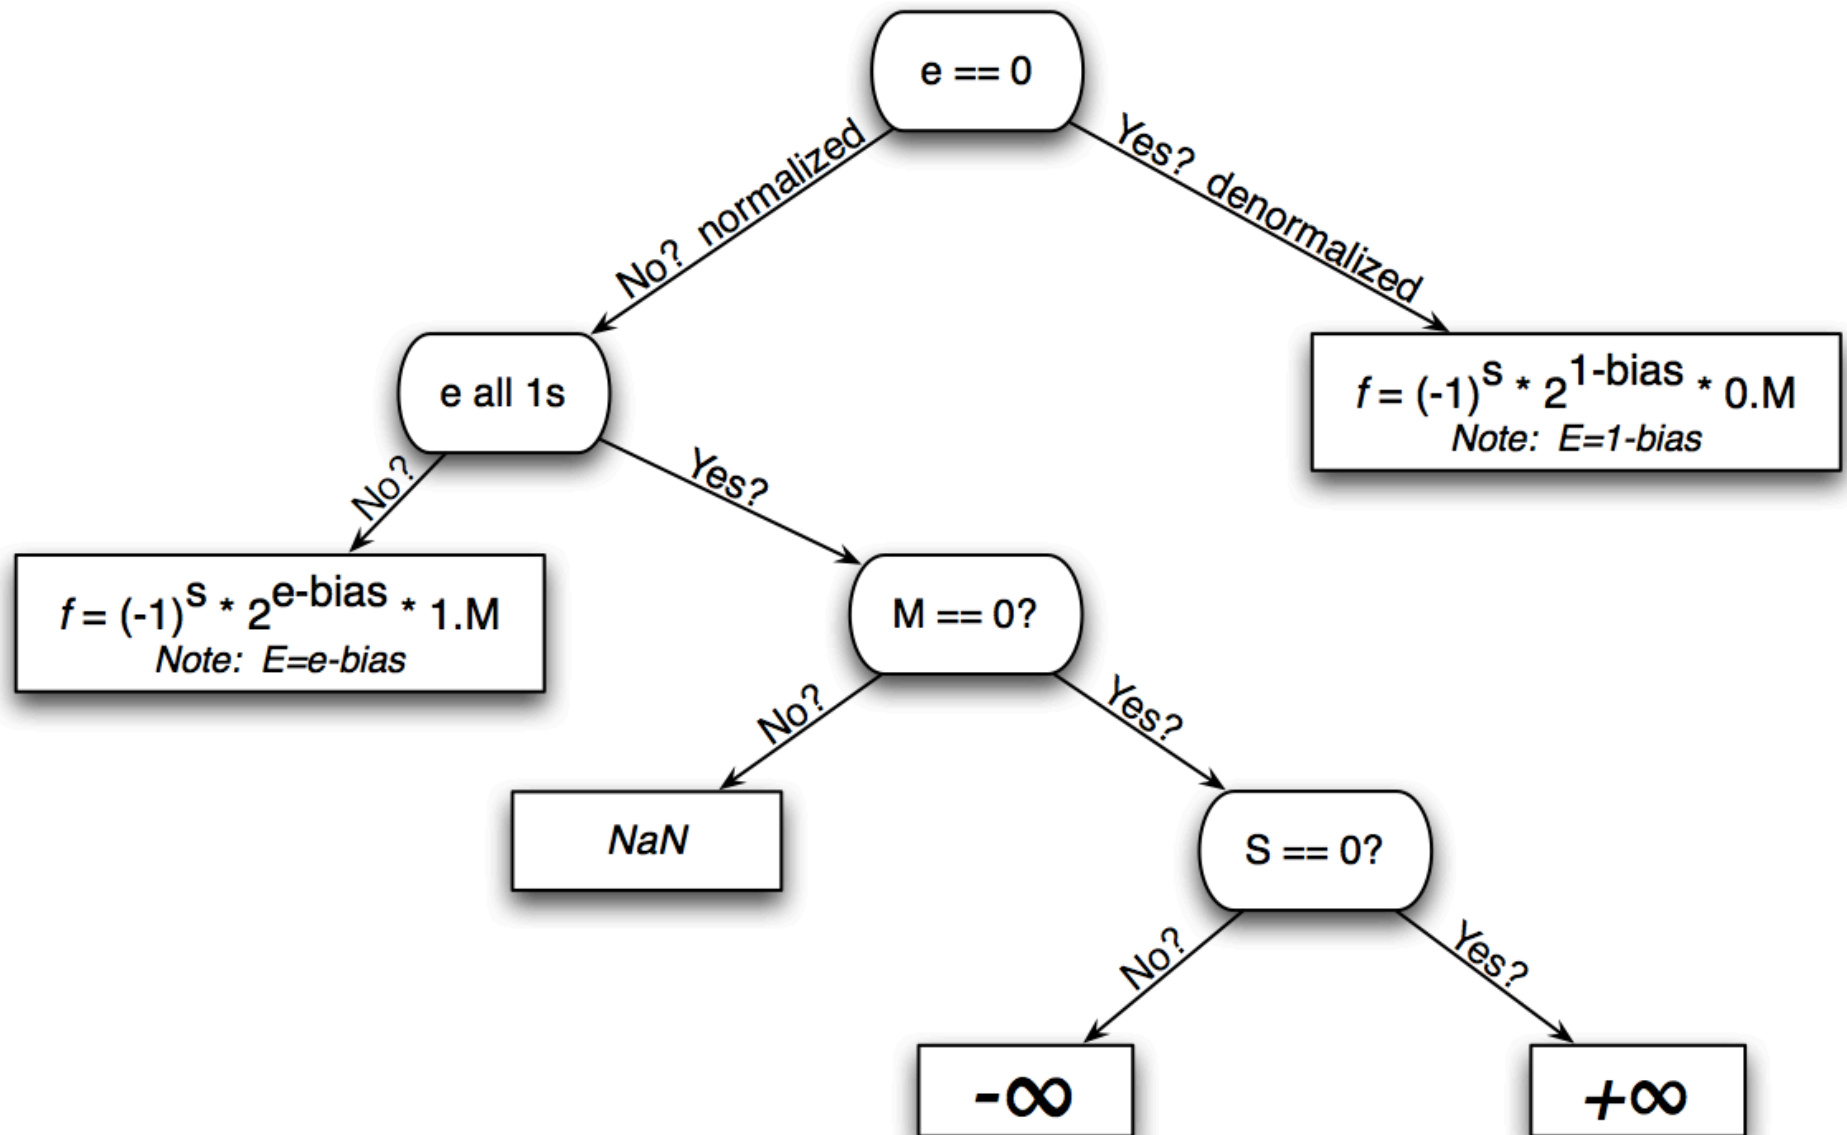
- With a 4-bit exponent, what will be the bias?
- Bias = $2^{\text{width}-1}-1 = 2^{4-1}-1 = 2^3-1 = 7$

Table of “Simple” Values

s	exp	frac	E	Value	comment
0	0000	000	-6	0	zero
0	0000	001	-6	$1/8 * 1/64 = 1/512$	closest to zero
0	0000	010	-6	$2/8 * 1/64 = 2/512$	
...	
0	0000	110	-6	$6/8 * 1/64 = 6/512$	
0	0000	111	-6	$7/8 * 1/64 = 7/512$	largest denormalized
0	0001	000	-6	$8/8 * 1/64 = 8/512$	smallest norm (rem. leading 1)
0	0001	001	-6	$9/8 * 1/64 = 9/512$	
...
0	0110	110	-1	$14/8 * 1/2 = 14/16$	
0	0110	111	-1	$15/8 * 1/2 = 15/16$	closest to 1 from below
0	0111	000	0	$8/8 * 1 = 1$	
0	0111	001	0	$9/8 * 1 = 9/8$	closest to 1 from above
0	0111	010	0	$10/8 * 1 = 10/8$	
...
0	1110	110	7	$14/8 * 128 = 224$	
0	1110	111	7	$15/8 * 128 = 240$	largest normalized
0	1111	000	n/a	infinity	

- Please convince yourselves that these make sense

Floating-point cheat sheet



“important” numbers

description	exp	frac
zero	00...00	00...00
smallest de-normalized	00...00	00...01
largest de-normalized	00...00	11...11
smallest normalized	00...01	00...00
one	01...11	00...00
largest normalized	11...10	11...11

Rounding

	\$1.40	\$1.60	\$1.50	\$2.50	-\$1.50
towards zero	\$1	\$1	\$1	\$2	-\$1
round down ($-\infty$)	\$1	\$1	\$1	\$2	-\$2
round up ($+\infty$)	\$2	\$2	\$2	\$3	-\$1
round to even	\$1	\$2	\$2	\$2	-\$2

Rounding

	\$1.40	\$1.60	\$1.50	\$2.50	-\$1.50
towards zero	\$1	\$1	\$1	\$2	-\$1
round down ($-\infty$)	\$1	\$1	\$1	\$2	-\$2
round up ($+\infty$)	\$2	\$2	\$2	\$3	-\$1
round to even	\$1	\$2	\$2	\$2	-\$2

- Don't be confused by round to even:
 - Round to the closest
 - Choose the even number when you're half-way between two possibilities

Why Round to Even?

- Rounding in same direction \rightarrow skew?
- Round to even:
 - sometimes up
 - sometimes down

The moral of the story ... back to our first example. Why false?

```
1  #include <stdio.h>
2
3  int main(int argc, char **argv)
4  {
5      float f=0.1;
6
7      if (f==0.1)
8          printf("It's 0.1\n");
9      else
10         printf("It's not 0.1\n");
11
12     return 0;
13 }
```

The moral of the story ... back to our first example. Why false?

```
1  #include <stdio.h>
2
3  int main(int argc, char **argv)
4  {
5      float f=0.1;
6
7      if (f==0.1)
8          printf("It's 0.1\n");
9      else
10         printf("It's not 0.1\n");
11
12     return 0;
13 }
```

- Can we store 0.1_{10} without rounding?
- Will there be less round error if we use a double?
- What happens if the compiler uses a double for the 0.1 in line 7? Will the values be the same?

These give us what we expect.

```
1  #include <stdio.h>
2
3  int main(int argc, char **argv)
4  {
5      float f = 0.1;
6
7      if (f==0.1f) {
8          printf("It's 0.1\n");
9      } else {
10         printf("It's not 0.1\n");
11     }
12
13     return 0;
14 }
```

```
1  #include <stdio.h>
2
3  int main(int argc, char **argv)
4  {
5      double f = 0.1;
6
7      if (f==0.1) {
8          printf("It's 0.1\n");
9      } else {
10         printf("It's not 0.1\n");
11     }
12
13     return 0;
14 }
```